



Pearson r vs. next-token entropy


Pythia-2.8B
(32 layers, $n=48$)




GPT-J-6B
(28 layers, $n=32$)



OPT-6.7B
(32 layers, $n=32$)



Pythia-6.9B
(32 layers, $n=32$)



GPT-2 XL
(48 layers, $n=48$)

